

State Farm Classification Coding Exercise

Part 4 - Build a Logistic Regression Model

A. Construct Logistic Regression Model Upon Training Data

Import numpy and pandas.

```
In [1]: import numpy as np  
import pandas as pd
```

Import data visualization libraries and set %matplotlib inline.

```
In [2]: import matplotlib.pyplot as plt  
import seaborn as sns  
%matplotlib inline
```

Import training model data into a Pandas dataframe called train_model1.

```
In [3]: train_model1 = pd.read_pickle('../State_Farm/Data/train_model.pickle')
```

Check number of rows and columns in train_model1 dataframe.

```
In [4]: train_model1.shape
```

```
Out[4]: (40000, 15)
```

View structure of train_model1 dataframe.

In [5]: `train_model1.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40000 entries, 0 to 39999
Data columns (total 15 columns):
y          40000 non-null int64
x75        40000 non-null float64
x37        40000 non-null float64
x58        40000 non-null float64
x97        40000 non-null float64
x41_flt    40000 non-null float64
x99        40000 non-null float64
x1         40000 non-null float64
x40        40000 non-null float64
x70        40000 non-null float64
x44        40000 non-null float64
x63        40000 non-null float64
x56        40000 non-null float64
x83        40000 non-null float64
x96        40000 non-null float64
dtypes: float64(14), int64(1)
memory usage: 4.6 MB
```

View first five rows of train_model1 dataframe.

In [6]: `train_model1.head()`

Out[6]:

	y	x75	x37	x58	x97	x41_flt	x99	x1	x40	
0	0	40.617107	-10.839200	2.078396	-2.125570	449.48	1.237667	74.425320	4.550518	41
1	1	-49.303165	57.917006	-2.696257	-36.030599	-525.06	1.952183	24.320711	-9.476135	36
2	1	-19.706659	-12.991058	-2.417447	26.212736	-599.50	0.558988	-66.160459	6.876065	49
3	0	-7.301283	37.658926	4.443710	19.221130	-220.71	0.214462	33.210943	-16.933984	82
4	0	-2.751656	-59.497091	-2.421952	-5.703269	-1405.59	-1.191319	-26.717872	-9.551514	-19

Gather summary statistics of train_model1 dataframe.

```
In [7]: train_model1.describe()
```

```
Out[7]:
```

	y	x75	x37	x58	x97	x41_fit	
count	40000.00000	40000.000000	40000.000000	40000.000000	40000.000000	40000.000000	40000.0
mean	0.20360	5.417519	4.814694	-0.892583	-2.514305	0.315315	0.0
std	0.40268	35.653631	31.561154	6.319470	18.551630	1001.659950	1.1
min	0.00000	-146.967384	-127.651997	-31.395877	-73.908741	-4496.460000	-4.3
25%	0.00000	-17.971049	-16.589036	-5.175537	-15.023076	-669.295000	-0.7
50%	0.00000	5.592658	4.493183	-0.891371	-2.514305	6.040000	0.0
75%	0.00000	28.988799	25.968530	3.413382	9.882079	678.325000	0.7
max	1.00000	144.548014	126.924294	22.420511	76.120119	4062.630000	4.4

Define X and y.

```
In [8]: feature_cols = ['x75', 'x37', 'x58', 'x97', 'x41_fit', 'x99', 'x1', 'x40', 'x70',
```

```
In [9]: X = train_model1[feature_cols]
y = train_model1['y']
print(X.shape)
print(y.shape)
```

```
(40000, 14)
(40000,)
```

Fit a logistic regression model on training model data.

```
In [10]: from sklearn.linear_model import LogisticRegression
from sklearn import metrics
```

```
In [11]: logreg = LogisticRegression(C=1e9)
logreg.fit(X, y)
```

```
Out[11]: LogisticRegression(C=1000000000.0, class_weight=None, dual=False,
fit_intercept=True, intercept_scaling=1, max_iter=100,
multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
solver='liblinear', tol=0.0001, verbose=0, warm_start=False)
```

Print logistic regression model intercept and coefficients.

```
In [12]: print(logreg.intercept_)
dict(zip(feature_cols, logreg.coef_[0]))

[-1.3847385]
```

```
Out[12]: {'x1': 0.007047416085280777,
'x37': -0.016105211331098775,
'x40': 0.00578891080883736,
'x41_flt': -0.0003401854661650888,
'x44': 0.006037683479860255,
'x56': -0.009542535312387216,
'x58': 0.11121565823866641,
'x63': 0.029668912026681572,
'x70': -0.005447882703786845,
'x75': -0.023153039385620668,
'x83': 0.08761590096736244,
'x96': -0.010344079456438539,
'x97': 0.03871572430063214,
'x99': 0.4788438298145828}
```

Express logistic regression model coefficients as odds.

```
In [13]: dict(zip(feature_cols, np.exp(logreg.coef_[0])))
```

```
Out[13]: {'x1': 1.007072307561193,
'x37': 0.9840237841566418,
'x40': 1.0058056989323634,
'x41_flt': 0.9996598723903498,
'x44': 1.0060559470287698,
'x56': 0.9905028501986991,
'x58': 1.1176359082580003,
'x63': 1.0301134193228654,
'x70': 0.9945669300975146,
'x75': 0.977112935566845,
'x83': 1.0915687709881947,
'x96': 0.9897092365401361,
'x97': 1.0394749441780715,
'x99': 1.6142070248626739}
```

Make predictions on training model data and calculate accuracy for training model data.

```
In [14]: y_pred_class = logreg.predict(X)
print(metrics.accuracy_score(y, y_pred_class).round(3))

0.843
```

Compute training model data null accuracy manually.

```
In [15]: print(1 - y.mean())

0.7964
```

Calculate area under the ROC curve (AUC) value for training data logistic regression model.


```
In [16]: y_pred_prob = logreg.predict_proba(X)[: , 1]
print(metrics.roc_auc_score(y, y_pred_prob).round(3))
```

0.854

Plot logistic regression model ROC curve for training data.

```
In [17]: thresholds = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9]
fpr, tpr, thresholds = metrics.roc_curve(y, y_pred_prob)
plt.plot(fpr, tpr)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
```

```
Out[17]: Text(0,0.5,'True Positive Rate (Sensitivity)')
```



Print confusion matrix to calculate training model data accuracy and error rates plus precision and recall.

```
In [18]: print(metrics.confusion_matrix(y, y_pred_class))
```

```
[[30265  1591]
 [ 4706  3438]]
```

Calculate training model data accuracy rate.

```
In [19]: float(30265 + 3438) / float(30265 + 1591 + 4706 + 3438)
```

```
Out[19]: 0.842575
```

Calculate training model data misclassification / error rate.

```
In [20]: float(4706 + 1591) / float(30265 + 1591 + 4706 + 3438)
```

```
Out[20]: 0.157425
```

Calculate precision to measure how confident the logistic regression model is for capturing the positives in training model data.

```
In [21]: float(3438) / float(3438 + 1591)
```

```
Out[21]: 0.6836349174786239
```

Calculate recall / sensitivity to measure how well the logistic regression model is capturing the positives in training model data.

```
In [22]: float(3438) / float(4706 + 3438)
```

```
Out[22]: 0.42215127701375244
```

Calculate specificity to measure how well the logistic regression model is capturing the negatives in training model data.

```
In [23]: float(30265) / float(30265 + 1591)
```

```
Out[23]: 0.9500565042692114
```

Print out training model data classification report for logistic regression model.

```
In [24]: from sklearn.metrics import classification_report
```

```
In [25]: print(classification_report(y, y_pred_class))
```

	precision	recall	f1-score	support
0	0.87	0.95	0.91	31856
1	0.68	0.42	0.52	8144
avg / total	0.83	0.84	0.83	40000

List out training model data logistic regression model accuracy scores and their average using 10-fold cross-validation.

In [26]: `from sklearn.cross_validation import cross_val_score`

C:\Users\kyrma\AppData\Local\Continuum\anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)

In [27]: `acc_scores = cross_val_score(logreg, X, y, cv=10, scoring='accuracy').round(3)`
`print(acc_scores)`
`print(acc_scores.mean().round(3))`

[0.837 0.839 0.843 0.84 0.842 0.847 0.84 0.846 0.849 0.844]
0.843

List out training model data logistic regression model AUC values and their average using 10-fold cross-validation.

In [28]: `auc_scores = cross_val_score(logreg, X, y, cv=10, scoring='roc_auc').round(3)`
`print(auc_scores)`
`print(auc_scores.mean().round(3))`

[0.849 0.845 0.849 0.854 0.85 0.86 0.853 0.854 0.87 0.858]
0.854

B. Generate Predictions in Cleaned Test Data

Import cleaned test data pickle file into a Pandas dataframe called test_model1.

In [29]: `test_model1 = pd.read_pickle('../State_Farm/Data/test_model.pickle')`

Check number of rows and columns in test_model1 dataframe.

In [30]: `test_model1.shape`

Out[30]: (10000, 14)

View structure of test_model1 dataframe.

In [31]: test_model1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
x75          10000 non-null float64
x37          10000 non-null float64
x58          10000 non-null float64
x97          10000 non-null float64
x41_flt      10000 non-null float64
x99          10000 non-null float64
x1           10000 non-null float64
x40          10000 non-null float64
x70          10000 non-null float64
x44          10000 non-null float64
x63          10000 non-null float64
x56          10000 non-null float64
x83          10000 non-null float64
x96          10000 non-null float64
dtypes: float64(14)
memory usage: 1.1 MB
```

View first five rows of test_model1 dataframe.

In [32]: test_model1.head()

Out[32]:

	x75	x37	x58	x97	x41_flt	x99	x1	x40	
0	80.851384	-32.086998	5.093232	13.739194	2475.46	1.141799	54.479467	-8.776231	44.269
1	-21.295879	29.391786	-5.989844	-13.018951	-1109.10	0.568757	-20.244923	23.337240	5.663
2	27.719905	-30.329997	6.115105	0.791425	-187.70	-0.816682	-61.467354	-8.845933	-47.647
3	-4.053955	11.088216	-2.750484	-16.716012	525.65	0.603007	-18.454831	-57.516611	-14.010
4	21.743536	-21.955105	-4.286183	23.003355	-1113.53	1.929231	15.810515	-13.207037	13.267

Define X_test.

In [33]: X_test = test_model1[feature_cols]
print(X_test.shape)

(10000, 14)

Make predictions on cleaned test data by calculating probabilities for belonging to positive class (labeled '1').

In [34]: y_test_pred_prob = logreg.predict_proba(X_test)[: , 1]

Construct results1_logreg Pandas dataframe out of test data y predicted class probabilities.


```
In [35]: results1_logreg = pd.DataFrame({'y_test_pred_prob': y_test_pred_prob[:]}))
results1_logreg.head()
```

Out[35]:

	y_test_pred_prob
0	0.146379
1	0.178329
2	0.211236
3	0.034441
4	0.413519

Check number of rows and columns in results1 dataframe.

```
In [36]: results1_logreg.shape
```

Out[36]: (10000, 1)

Export results1 Pandas dataframe containing test data y predicted class probabilities to CSV file.

```
In [37]: results1_logreg.to_csv('../State_Farm/Data/results1_logreg.csv', sep=',', index=False)
```